

Fossil Identification App - Full Stack Roadmap

1. Architecture Overview

The Fossil Identification App integrates Flutter (frontend), FastAPI (backend), Supabase (auth, storage, database), and an AI model (Claude/GPT-4 Vision) for fossil recognition using a structured rule-based system inspired by a 66-page document.

2. Should You Use RAG?

No. RAG is unnecessary because the fossil identification document is structured and rule-based, not a knowledge base.

Instead, pass the document as a **system prompt** to the AI model.

3. Backend Choice: FastAPI ■

- Async, lightweight, and easy to scale.
- Handles file uploads and integrates well with Supabase.
- Generates automatic Swagger documentation.

Alternative: Express.js with TypeScript if you prefer Node, but FastAPI is better for AI workloads.

4. Core Architecture Layout

project/

■ ■ ■ app/

■ ■ ■ main.py

■ ■ ■ routers/

```

■ ■ ■ ■ ■ fossils.py

■ ■ ■ ■ ■ services/

■ ■ ■ ■ ■ ai_service.py

■ ■ ■ ■ ■ storage.py

■ ■ ■ ■ ■ models/

■ ■ ■ ■ ■ schemas.py

■ ■ ■ ■ ■ prompts/

■ ■ ■ ■ ■ fossil_prompt.txt

■ ■ ■ ■ ■ .env

---

```

5. Backend Essentials

Security:

- Use API keys from environment variables.
- Enable JWT/Supabase Auth.
- Rate-limit requests (e.g., 10/min per user).
- Restrict CORS to your Flutter app domain.
- Validate uploads (JPG/PNG only, <10MB).

Critical Tasks:

Priority	Task	Description
■	Image Preprocessing	Resize/compress before sending to AI
■	Error Handling	User-friendly messages for AI/DB failures
■	Caching	Prevent duplicate analyses
■	Logging	Track API requests and AI cost
■	Monitoring	Sentry or similar

6. Supabase Integration

- Auth: Secure login with email/password or OAuth.
- Storage: Fossil images.
- Database: User data, history, AI results.
- Optional: Realtime updates via Supabase Realtime.

7. AI Integration (Claude/OpenAI)

****No RAG needed.****

- Load fossil prompt as `fossil_prompt.txt`
- Send images and metadata together.
- Expect structured JSON output from AI.

```
```python
```

```
class FossilAI:
```

```
 async def analyze_fossil(self, image_bytes, metadata):
```

```
 response = await client.messages.create(
```

```
 model="claude-sonnet-4",
```

```
 max_tokens=4096,
```

```
 messages=[
```

```
 {"role": "user",
```

```
 "content": [
```

```
 {"type": "image", "source": {"type": "base64", "data": base64.b64encode(image_bytes).decode()}},
```

```
 {"type": "text", "text": f"{self.prompt}\n\nMetadata: {metadata}"}]
```

```
]
```

```
]]
```

)

```
return response.content[0].text
```

...

---

## 8. Rule Engine Overview

Implements the **4-step fossil classification funnel** from the document:

1. Coarse Category → (tooth, bone, shell, plant, etc.)
2. Environment → (marine, terrestrial, freshwater)
3. Family Candidates → based on diagnostic rules
4. Species/Genus → only when diagnostic cues are clear

This gives deterministic, testable results.

---

## 9. Production Deployment Checklist

- HTTPS enabled (Railway, Render, or Cloud Run)
- Environment variables set (.env)
- Health endpoint (/health)
- Logging and rate limiting
- Sentry for errors

---

## 10. Cost Optimization

- Cache results by image hash.
- Resize images to max 1024×1024.
- Limit 5 analyses/min per user.

- Store repeated results in Supabase.

---

## 11. Phase-wise Implementation Plan

### ***Phase 1: Core Setup (Weeks 1–2)***

- Set up Supabase (auth, DB, storage)
- Implement FastAPI skeleton & image upload
- Integrate Claude/OpenAI API

### ***Phase 2: Rule Engine (Weeks 3–4)***

- Implement CoarseCategory & Environment logic
- Add whitelist for families
- Validate JSON schema outputs

### ***Phase 3: Production & Optimization (Weeks 5–6)***

- Add caching, error handling
- Optimize model prompts
- Deploy to Railway/Render
- Add monitoring & logs

---

## 12. Success Metrics

- Accuracy: >90% on test fossils
- Response time: <2s
- Cost: <\$0.02 per image
- Reliability: 99.9% uptime